

6502 Instruction Set



TOC: [Description](#) / [Instructions by Type](#) / [Address Modes in Detail](#) / [Instructions in Detail](#) / ["Illegal" Opcodes](#) / [Comparisons](#) /

[A Primer of 6502 Arithmetic Operations](#) / [Jump Vectors and Stack Operations](#) / [Instruction Layout](#) / [Pinout](#) / [65xx-Family](#)

Tools: [6502 Emulator](#) / [6502 Assembler](#) / [6502 Disassembler](#)

HI	LO-NIBBLE															
	-0	-1	-2	-3	-4	-5	-6	-7	-8	-9	-A	-B	-C	-D	-E	-F
0-	BRK impl	ORA X,ind				ORA zpg	ASL zpg		PHP impl	ORA #	ASL A			ORA abs	ASL abs	
1-	BPL rel	ORA ind,Y				ORA zpg,X	ASL zpg,X		CLC impl	ORA abs,Y				ORA abs,X	ASL abs,X	
2-	JSR abs	AND X,ind			BIT zpg	AND zpg	ROL zpg		PLP impl	AND #	ROL A		BIT abs	AND abs	ROL abs	
3-	BMI rel	AND ind,Y				AND zpg,X	ROL zpg,X		SEC impl	AND abs,Y				AND abs,X	ROL abs,X	
4-	RTI impl	EOR X,ind				EOR zpg	LSR zpg		PHA impl	EOR #	LSR A		JMP abs	EOR abs	LSR abs	
5-	BVC rel	EOR ind,Y				EOR zpg,X	LSR zpg,X		CLI impl	EOR abs,Y				EOR abs,X	LSR abs,X	
6-	RTS impl	ADC X,ind				ADC zpg	ROR zpg		PLA impl	ADC #	ROR A		JMP ind	ADC abs	ROR abs	
7-	BVS rel	ADC ind,Y				ADC zpg,X	ROR zpg,X		SEI impl	ADC abs,Y				ADC abs,X	ROR abs,X	
8-		STA X,ind			STY zpg	STA zpg	STX zpg		DEY impl		TXA impl		STY abs	STA abs	STX abs	
9-	BCC rel	STA ind,Y			STY zpg,X	STA zpg,X	STX zpg,Y		TYA impl	STA abs,Y	TXS impl			STA abs,X		
A-	LDY #	LDA X,ind	LDX #		LDY zpg	LDA zpg	LDX zpg		TAY impl	LDA #	TAX impl		LDY abs	LDA abs	LDX abs	
B-	BCS rel	LDA ind,Y			LDY zpg,X	LDA zpg,X	LDX zpg,Y		CLV impl	LDA abs,Y	TSX impl		LDY abs,X	LDA abs,X	LDX abs,Y	
C-	CPY #	CMP X,ind			CPY zpg	CMP zpg	DEC zpg		INY impl	CMP #	DEX impl		CPY abs	CMP abs	DEC abs	
D-	BNE rel	CMP ind,Y				CMP zpg,X	DEC zpg,X		CLD impl	CMP abs,Y				CMP abs,X	DEC abs,X	
E-	CPX #	SBC X,ind			CPX zpg	SBC zpg	INC zpg		INX impl	SBC #	NOP impl		CPX abs	SBC abs	INC abs	
F-	BEQ rel	SBC ind,Y				SBC zpg,X	INC zpg,X		SED impl	SBC abs,Y				SBC abs,X	INC abs,X	

☐ show illegal opcodes

Description

Address Modes

A Accumulator	OPC A	<i>op is AC (implied single byte instruction)</i>
abs absolute	OPC \$LLHH	<i>op is address \$HHLL *</i>
abs,X absolute, X-indexed	OPC \$LLHH,X	<i>op is address; effective address is address incremented by X with carry **</i>
abs,Y absolute, Y-indexed	OPC \$LLHH,Y	<i>op is address; effective address is address incremented by Y with carry **</i>
# immediate	OPC #\$BB	<i>op is byte BB</i>
impl implied	OPC	<i>op implied</i>
ind indirect	OPC (\$LLHH)	<i>op is address; effective address is contents of word at address: C.w(\$HHLL)</i>
X,ind X-indexed, indirect	OPC (\$LL,X)	<i>op is zp address; effective address is word in (LL + X, LL + X + 1), inc. w/o carry: C.w(\$00LL + X)</i>
ind,Y indirect, Y-indexed	OPC (\$LL),Y	<i>op is zp address; effective address is word in (LL, LL + 1) incremented by Y with carry: C.w(\$00LL) + Y</i>
rel relative	OPC \$BB	<i>branch target is PC + signed offset BB ***</i>
zpg zp	OPC \$LL	<i>op is zp address (hi-byte is zero, address = \$00LL)</i>
zpg,X zp, X-indexed	OPC \$LL,X	<i>op is zp address; effective address is address incremented by X w/o carry **</i>
zpg,Y zp, Y-indexed	OPC \$LL,Y	<i>op is zp address; effective address is address incremented by Y w/o carry **</i>

* 16-bit address words are little endian, lo(w)-byte first, followed by the hi(gh)-byte.
(An assembler will use a human readable, big-endian notation as in \$HHLL.)

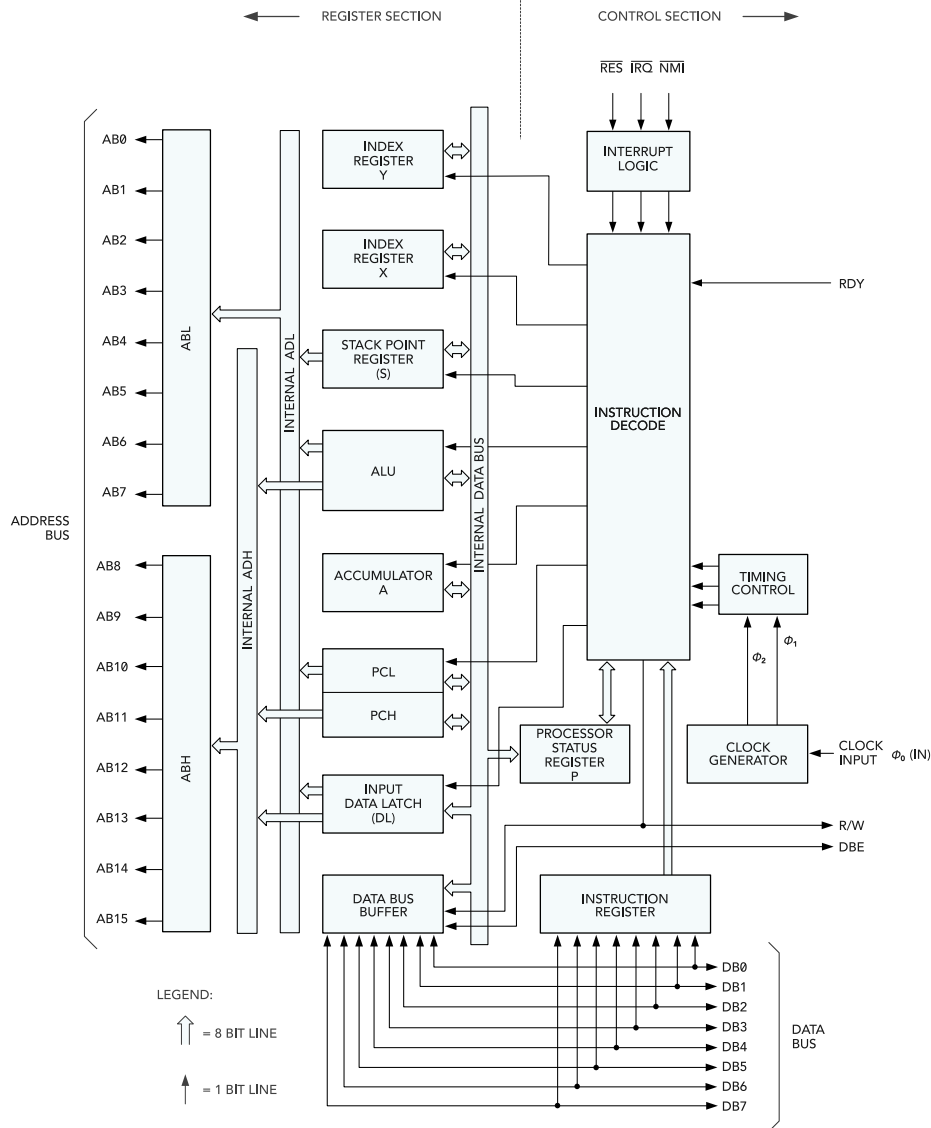
** The available 16-bit address space is conceived as consisting of pages of 256 bytes each, with address hi-bytes representing the page index. An increment with carry may affect the hi-byte and may thus result in a crossing of page boundaries, adding an extra cycle to the execution. Increments without carry do not affect the hi-byte of an address and no page transitions do occur. Generally, increments of 16-bit addresses include a carry, increments of zeropage addresses don't. Notably this is not related in any way to the state of the carry bit of the accumulator.

*** Branch offsets are signed 8-bit values, -128 ... +127, negative offsets in two's complement. Page transitions may occur and add an extra cycle to the execution.

Instructions by Name

[ADC](#) add with carry

[AND](#) and (with accumulator)
[ASL](#) arithmetic shift left
[BCC](#) branch on carry clear
[BCS](#) branch on carry set
[BEQ](#) branch on equal (zero set)
[BIT](#) bit test
[BMI](#) branch on minus (negative set)
[BNE](#) branch on not equal (zero clear)
[BPL](#) branch on plus (negative clear)
[BRK](#) break / interrupt
[BVC](#) branch on overflow clear
[BVS](#) branch on overflow set
[CLC](#) clear carry
[CLD](#) clear decimal
[CLI](#) clear interrupt disable
[CLV](#) clear overflow
[CMP](#) compare (with accumulator)
[CPX](#) compare with X
[CPY](#) compare with Y
[DEC](#) decrement
[DEX](#) decrement X
[DEY](#) decrement Y
[EOR](#) exclusive or (with accumulator)
[INC](#) increment
[INX](#) increment X
[INY](#) increment Y
[JMP](#) jump
[JSR](#) jump subroutine
[LDA](#) load accumulator
[LDX](#) load X
[LDY](#) load Y
[LSR](#) logical shift right
[NOP](#) no operation
[ORA](#) or with accumulator
[PHA](#) push accumulator
[PHP](#) push processor status (SR)
[PLA](#) pull accumulator
[PLP](#) pull processor status (SR)
[ROL](#) rotate left
[ROR](#) rotate right
[RTI](#) return from interrupt
[RTS](#) return from subroutine
[SBC](#) subtract with carry
[SEC](#) set carry
[SED](#) set decimal
[SEI](#) set interrupt disable
[STA](#) store accumulator
[STX](#) store X
[STY](#) store Y
[TAX](#) transfer accumulator to X
[TAY](#) transfer accumulator to Y
[TSX](#) transfer stack pointer to X
[TXA](#) transfer X to accumulator
[TXS](#) transfer X to stack pointer
[TYA](#) transfer Y to accumulator



Block diagram of the NMOS 6502 CPU.

After (6501 specifics omitted):

MCS6502 Microcomputer Family Hardware Manual; January 1976.

MOS Technology, Inc., Norristown/PA, 1976.

Registers

PC	program counter	(16 bit)
AC	accumulator	(8 bit)
X	X register	(8 bit)
Y	Y register	(8 bit)
SR	status register [NV-BDIZC]	(8 bit)
SP	stack pointer	(8 bit)

Note: The status register (SR) is also known as the P register, the accumulator (AC) as just A, the stack pointer (SP) as S, and X and Y registers as XR and YR respectively.

Machine language monitors show them typically like,

```
PC  IRQ  SR  AC  XR  YR  SP
0401 E62E 32 04 5E 00 F8
```

("IRQ" is not a register, but the interrupt request vector, see below.)

- The **accumulator** is the main register of the 6502. Its content is typically used by the arithmetic logic unit (ALU) for the first operand and results are deposited in the accumulator again. Thus its name, as results accumulate in this register. Most arithmetic and logical operations interact with this register.
- The **X** and **Y** registers are auxiliary registers. Like the accumulator, they can be loaded directly with values, both

immediately (as literal constants) or from memory. Additionally, they can be incremented and decremented, and their contents may be transferred to and from the accumulator. Their main purpose is the use as index registers, where their contents is added to a base memory location, before any values are either stored to or retrieved from the resulting address, which is known as the *effective address*. This is commonly used for loops and table lookups at a given index, hence the name. (See [address modes](#), below.)

- The **program counter** keeps track of the memory location holding the current instruction code. Its contents is automatically stepped up as the program is executed and is modified by branch and jump operations. As it must be able to address the full 16-bit address range of 64K bytes, it's the only 16-bit register of the 6502.
- The **stack pointer** points to the current *top of stack* (or rather, to its bottom, as the stack grows top-down.) The processor stack is located on memory page #1 (\$0100-\$01FF), a 256 bytes *last-in-first-out* (LIFO) stack, which enables subroutines and also serves as a quick intermediate storage. As a 8-bit register, the stack pointer holds just the low-byte of this address (the offset from \$0100.) Be aware that this will just wrap around, in case that the stack underflows.
- The **status register** holds the status of the processor, consisting of flags reflecting results of previous operations, configuration flags, like disabeling (blocking) interrupts or setting up binary encoded decimal mode (BCD), and the carry flag, which enables multi-byte arithmetics.

Status Register Flags (bit 7 to bit 0)

N Negative
V Overflow
- ignored
B Break
D Decimal (use BCD for arithmetics)
I Interrupt (IRQ disable)
Z Zero
C Carry

- The **zero flag** (Z) indicates a value of all zero bits and the **negative flag** (N) indicates the presence of a set sign bit in bit-position 7. These flags are always updated, whenever a value is transferred to a CPU register (A,X,Y) and as a result of any logical ALU operations. The Z and N flags are also updated by increment and decrement operations acting on a memory location.
- The **carry flag** (C) flag is used as a buffer and as a borrow in arithmetic operations. Any comparisons will update this additionally to the Z and N flags, as do shift and rotate operations.
- All arithmetic operations update the Z, N, C and V flags.
- The **overflow flag** (V) indicates overflow with signed binary arithmetics. As a signed byte represents a range of -128 to +127, an overflow can never occur when the operands are of opposite sign, since the result will never exceed this range. Thus, overflow may only occur, if both operands are of the same sign. Then, the result must be also of the same sign. Otherwise, overflow is detected and the overflow flag is set. (I.e., both operands have a zero in the sign position at bit 7, but bit 7 of the result is 1, or, both operands have the sign-bit set, but the result is positive.)
- The **decimal flag** (D) sets the ALU to binary coded decimal (BCD) mode for additions and subtractions (ADC, SBC).
- The **interrupt inhibit flag** (I) blocks any maskable interrupt requests (IRQ).

- The **break flag** (B) is not an actual flag implemented in a register, and rather appears only, when the status register is pushed onto or pulled from the stack. When pushed, it will be 1 when transferred by a BRK or PHP instruction, and zero otherwise (i.e., when pushed by a hardware interrupt). When pulled into the status register (by PLP or on RTI), it will be ignored.
In other words, the break flag will be inserted, whenever the status register is transferred to the stack by software (BRK or PHP), and will be zero, when transferred by hardware. Since there is no actual slot for the break flag, it will be always ignored, when retrieved (PLP or RTI). The break flag is not accessed by the CPU at anytime and there is no internal representation. Its purpose is more for patching, to discern an interrupt caused by a BRK instruction from a normal interrupt initiated by hardware.
- Any of these flags (but the break flag) may be set or cleared by dedicated instructions. Moreover, there are branch instructions to conditionally divert the control flow depending on the respective state of the Z, N, C or V flag.

Processor Stack

LIFO, top-down, 8 bit range, 0x0100 - 0x01FF

Bytes, Words, Addressing

8 bit bytes, 16 bit words in lobyte-hibyte representation (Little-Endian).
16 bit address range, operands follow instruction codes.

Signed values are two's complement, sign in bit 7 (most significant bit).
(%11111111 = \$FF = -1, %10000000 = \$80 = -128, %01111111 = \$7F = +127)
Signed binary and binary coded decimal (BCD) arithmetic modes.

System Vectors

\$FFFA, \$FFFB ... NMI (Non-Maskable Interrupt) vector, 16-bit (LB, HB)
\$FFFC, \$FFFD ... RES (Reset) vector, 16-bit (LB, HB)
\$FFFE, \$FFFF ... IRQ (Interrupt Request) vector, 16-bit (LB, HB)

Start/Reset Operations

An active-low reset line allows to hold the processor in a known disabled state, while the system is initialized. As the reset line goes high, the processor performs a start sequence of 7 cycles, at the end of which the program counter (PC) is read from the address provided in the 16-bit reset vector at \$FFFC (LB-HB). Then, at the eighth cycle, the processor transfers control by performing a JMP to the provided address.

Any other initializations are left to the thus executed program. (Notably, instructions exist for the initialization and loading of all registers, but for the program counter, which is provided by the reset vector at \$FFFC.)

Instructions by Type

• Transfer Instructions

Load, store, interregister transfer

[LDA](#) load accumulator

[LDX](#) load X

[LDY](#) load Y

[STA](#) store accumulator

[STX](#) store X

[STY](#) store Y

[TAX](#) transfer accumulator to X

[TAY](#) transfer accumulator to Y

[TSX](#) transfer stack pointer to X
[TXA](#) transfer X to accumulator
[TXS](#) transfer X to stack pointer
[TYA](#) transfer Y to accumulator

• Stack Instructions

These instructions transfer the accumulator or status register (flags) to and from the stack. The processor stack is a last-in-first-out (LIFO) stack of 256 bytes length, implemented at addresses \$0100 - \$01FF. The stack grows down as new values are pushed onto it with the current insertion point maintained in the stack pointer register.

(When a byte is pushed onto the stack, it will be stored in the address indicated by the value currently in the stack pointer, which will be then decremented by 1. Conversely, when a value is pulled from the stack, the stack pointer is incremented. The stack pointer is accessible by the [TSX](#) and [TXS](#) instructions.)

[PHA](#) push accumulator
[PHP](#) push processor status register (with break flag set)
[PLA](#) pull accumulator
[PLP](#) pull processor status register

• Decrements & Increments

[DEC](#) decrement (memory)
[DEX](#) decrement X
[DEY](#) decrement Y
[INC](#) increment (memory)
[INX](#) increment X
[INY](#) increment Y

• Arithmetic Operations

[ADC](#) add with carry (prepare by [CLC](#))
[SBC](#) subtract with carry (prepare by [SEC](#))

See the [Primer of 6502 Arithmetic Instructions](#) below for details.

• Logical Operations

[AND](#) and (with accumulator)
[EOR](#) exclusive or (with accumulator)
[ORA](#) (inclusive) or with accumulator

• Shift & Rotate Instructions

All shift and rotate instructions preserve the bit shifted out in the carry flag.

[ASL](#) arithmetic shift left (shifts in a zero bit on the right)
[LSR](#) logical shift right (shifts in a zero bit on the left)
[ROL](#) rotate left (shifts in carry bit on the right)
[ROR](#) rotate right (shifts in zero bit on the left)

• Flag Instructions

[CLC](#) clear carry
[CLD](#) clear decimal (BCD arithmetics disabled)
[CLI](#) clear interrupt disable
[CLV](#) clear overflow
[SEC](#) set carry
[SED](#) set decimal (BCD arithmetics enabled)
[SEI](#) set interrupt disable

• Comparisons

Generally, comparison instructions subtract the operand from the given register without affecting this register. Flags are still set

as with a normal subtraction and thus the relation of the two values becomes accessible by the Zero, Carry and Negative flags.
(See the branch instructions below for how to evaluate flags.)

Relation <i>R - Op</i>	Z	C	N
Register < Operand	0	0	<i>sign bit of result</i>
Register = Operand	1	1	0
Register > Operand	0	1	<i>sign bit of result</i>

[CMP](#) compare (with accumulator)

[CPX](#) compare with X

[CPY](#) compare with Y

- **Conditional Branch Instructions**

Branch targets are relative, signed 8-bit address offsets. (An offset of #0 corresponds to the immediately following address – or a rather odd and expensive NOP.)

[BCC](#) branch on carry clear

[BCS](#) branch on carry set

[BEQ](#) branch on equal (zero set)

[BMI](#) branch on minus (negative set)

[BNE](#) branch on not equal (zero clear)

[BPL](#) branch on plus (negative clear)

[BVC](#) branch on overflow clear

[BVS](#) branch on overflow set

- **Jumps & Subroutines**

JSR and RTS affect the stack as the return address is pushed onto or pulled from the stack, respectively.

(JSR will first push the high-byte of the return address [PC+2] onto the stack, then the low-byte. The stack will then contain, seen from the bottom or from the most recently added byte, [PC+2]-L [PC+2]-H.)

[JMP](#) jump

[JSR](#) jump subroutine

[RTS](#) return from subroutine

- **Interrupts**

A hardware interrupt (maskable IRQ and non-maskable NMI), will cause the processor to put first the address currently in the program counter onto the stack (in HB-LB order), followed by the value of the status register. (The stack will now contain, seen from the bottom or from the most recently added byte, SR PC-L PC-H with the stack pointer pointing to the address below the stored contents of status register.) Then, the processor will divert its control flow to the address provided in the two word-size interrupt vectors at \$FFFA (IRQ) and \$FFFE (NMI).

A set interrupt disable flag will inhibit the execution of an IRQ, but not of a NMI, which will be executed anyways.

The break instruction (BRK) behaves like a NMI, but will push the value of PC+2 onto the stack to be used as the return address. Also, as with any software initiated transfer of the status register to the stack, the break flag will be found set on the respective value pushed onto the stack. Then, control is transferred to the address in the NMI-vector at \$FFFE.

In any way, the interrupt disable flag is set to inhibit any further IRQ as control is transferred to the interrupt handler specified by the respective interrupt vector.

The RTI instruction restores the status register from the stack and behaves otherwise like the JSR instruction. (The break flag is always ignored as the status is read from the stack, as it isn't a real processor flag anyway.)

[BRK](#) break / software interrupt

[RTI](#) return from interrupt

- **Other**

BIT bit test (accumulator & memory)

NOP no operation

6502 Address Modes in Detail

(This section, especially the diagrams included, is heavily inspired by the Acorn Atom manual "[Atomic Theory and Practice](#)" by David Johnson Davies, Acorn Computers Limited, 2nd ed. 1980, p 118-121.)

- **Implied Addressing**

These instructions act directly on one or more registers or flags internal to the CPU. Therefore, these instructions are principally single-byte instructions, lacking an explicit operand. The operand is implied, as it is already provided by the very instruction.

Instructions targeting exclusively the contents of the accumulator may or may not be denoted by using an explicit "A" as the operand, depending on the flavor of syntax. (This may be regarded as a special address mode of its own, but it is really a special case of an implied instruction. It is still a single-byte instruction and no operand is provided in machine language.)

Mnemonic Examples:

CLCclear the carry flag
 ROL Arotate contents of accumulator left by one position
 ROLsame as above, implicit notation (A implied)
 TXAtransfer contents of X-register to the accumulator
 PHApush the contents of the accumulator to the stack
 RTSreturn from subroutine (by pulling PC from stack)

Mind that some of these instructions, while simple in appearance, may be quite complex operations, like "PHA", which involves the accumulator, the stack pointer and memory access.

- **Immediate Addressing**

Here, a literal operand is given immediately after the instruction. The operand is always an 8-bit value and the total instruction length is always 2 bytes. In memory, the operand is a single byte following immediately after the instruction code. In assembler, the mode is usually indicated by a "#" prefix adjacent to the operand.

Mnemonic Instruction

LDA #7

A9	07
----	----

 
 A:

07


Mnemonic Examples:

LDA #\$07load the literal hexadecimal value "\$7" into the accumulator
 ADC #\$A0add the literal hexadecimal value "\$A0" to the accumulator
 CPX #\$32compare the X-register to the literal hexadecimal value "\$32"

- **Absolute Addressing**

Absolute addressing modes provides the 16-bit address of a memory location, the contents of which used as the operand to the instruction. In machine language, the address is provided in two bytes immediately after the instruction (making these 3-byte instructions) in low-byte, high-byte order (LLHH) or little-endian. In assembler, conventional numbers (HHLL order or big-endian words) are used to provide the address.

Absolute addresses are also used for the jump instructions JMP and JSR to provide the address for the next instruction to continue with in the control flow.




Mnemonic Examples:

LDA \$3010load.the.contents of address "\$3010" into the accumulator
 ROL \$08A0rotate.the contents of address "\$08A0" left by one position
 JMP \$4000jump.to.(continue with) location "\$4000"

• Zero-Page Addressing

The 16-bit address space available to the 6502 is thought to consist of 256 "pages" of 256 memory locations each (\$00...\$FF). In this model the high-byte of an address gives the page number and the low-byte a location inside this page. The very first of these pages, where the high-byte is zero (addresses \$0000...\$00FF), is somewhat special.

The zero-page address mode is similar to absolute address mode, but these instructions use only a single byte for the operand, the low-byte, while the high-byte is assumed to be zero by definition. Therefore, these instructions have a total length of just two bytes (one less than absolute mode) and take one CPU cycle less to execute, as there is one byte less to fetch.



Mnemonic Examples:


LDA \$80load.the contents of address "\$0080" into the accumulator
 BIT \$A2perform bit-test with the contents of address "\$00A2"
 ASL \$9Aarithmetic shift left of the contents of location "\$009A"

(One way to think of the zero-page is as a page of 256 additional registers, somewhat slower than the internal registers, but with zero-page instructions also faster executing than "normal" instructions. The zero-page has a few more tricks up its sleeve, making these addresses perform more like real registers, see below.)

• Indexed Addressing: Absolute,X and Absolute,Y

Indexed addressing adds the contents of either the X-register or the Y-register to the provided address to give the *effective address*, which provides the operand.

These instructions are useful to e.g., load values from tables or to write to a continuous segment of memory in a loop. The most basic forms are "absolute,X" and "absolute,Y", where either the X- or the Y-register, respectively, is added to a given base address. As the base address is a 16-bit value, these are generally 3-byte instructions. Since there is an additional operation to perform to determine the effective address, these instructions are one cycle slower than those using absolute addressing mode.*



Mnemonic Examples:

LDA \$3120,Xload.the.contents of address "\$3120 + X" into A
 LDX \$8240,Yload.the.contents of address "\$8240 + Y" into X
 INC \$1400,Xincrement the contents of address "\$1400 + X"


*) If the addition of the contents of the index register effects in a change of the high-byte given by the base address so that the effective address is on the next memory page, the additional operation to increment the high-byte takes another CPU cycle. This is also known as a crossing of page boundaries.

• Indexed Addressing: Zero-Page,X (and Zero-Page,Y)

As with absolute addressing, there is also a zero-page mode for indexed addressing. However, this is generally only available with the X-register. (The only exception to this is LDX, which has an indexed zero-page mode utilizing the Y-register.)

As we have already seen with normal zero-page mode, these instructions are one byte less in total length (two bytes) and take one CPU cycle less than instructions in absolute indexed mode.

Unlike absolute indexed instructions with 16-bit base addresses, zero-page indexed instructions never affect the high-byte of the effective address, which will simply wrap around in the zero-page, and there is no penalty for crossing any page boundaries.



Mnemonic Examples:


LDA \$80,Xload.the.contents of address "\$0080 + X" into A
 LSR \$82,Xshift.the contents of address "\$0082 + X" left
 LDX \$60,Yload.the.contents of address "\$0060 + Y" into X

• Indirect Addressing

This mode looks up a given address and uses the contents of this address and the next one (in LLHH little-endian order) as the effective address. In its basic form, this mode is available for the JMP instruction only. (Its generally use is jump vectors and jump tables.)

Like the absolute JMP instruction it uses a 16-bit address (3 bytes in total), but takes two additional CPU cycles to execute, since there are two additional bytes to fetch for the lookup of the effective jump target.

Generally, indirect addressing is denoted by putting the lookup address in parenthesis.



Mnemonic Example:

JMP (\$FF82) ...jump.to.address given in addresses "\$FF82" and "\$FF83"


• Pre-Indexed Indirect, "(Zero-Page,X)"

Indexed indirect address modes are generally available only for instructions supplying an operand to the accumulator (LDA, STA, ADC, SBC, AND, ORA, EOR, etc). The placement of the index register inside or outside of the parenthesis indicating the address lookup will give you clue what these instructions are doing.

Pre-indexed indirect address mode is only available in combination with the X-register. It works much like the "zero-page,X" mode, but, after the X-register has been added to the base address, instead of directly accessing this, an additional lookup is performed, reading the contents of resulting address and the next one (in LLHH little-endian order), in order to determine the effective address.

Like with "zero-page,X" mode, the total instruction length is 2 bytes, but there are two additional CPU cycles in order to fetch the effective 16-bit address. As "zero-page,X" mode, a lookup address will never overflow into the next page, but will simply wrap around in the zero-page.

These instructions are useful, whenever we want to loop over a table of pointers to disperse addresses, or where we want to apply the same operation to various addresses, which we have stored as a table in the zero-page.



Mnemonic Examples:

LDA (\$70,X) ...load the contents of the location given in addresses "\$0070+X" and "\$0070+1+X" into A

STA (\$A2,X) ...store the contents of A in the location given in addresses "\$00A2+X" and "\$00A3+X"


EOR (\$BA,X) ...perform an exclusive OR of the contents of A and the contents of the location given in addresses "\$00BA+X" and "\$00BB+X"

• Post-Indexed Indirect, "(Zero-Page),Y"

Post-indexed indirect addressing is only available in combination with the Y-register. As indicated by the indexing term ",Y" being appended to the outside of the parenthesis indicating the indirect lookup, here, a pointer is first read (from the given zero-page address) and resolved and only then the contents of the Y-register is added to this to give the effective address.

Like with "zero-page,Y" mode, the total instruction length is 2 bytes, but there it takes an additional CPU cycles to resolve and index the 16-bit pointer. As with "absolute,X" mode, the effective address may overflow into the next page, in the case of which the execution uses an extra CPU cycle.

These instructions are useful, wherever we want to perform lookups on varying bases addresses or whenever we want to loop over tables, the base address of which we have stored in the zero-page.



Mnemonic Examples:

LDA (\$70),Y ...add the contents of the Y-register to the pointer provided in "\$0070" and "\$0071" and load the contents of this address into A

STA (\$A2),Y ...store the contents of A in the location given by the pointer in "\$00A2" and "\$00A3" plus the contents of the Y-register

EOR (\$BA),Y ...perform an exclusive OR of the contents of A and the address given by the addition of Y to the pointer in "\$00BA" and "\$00BB"

• Relative Addressing (Conditional Branching)

This final address mode is exclusive to conditional branch instructions, which branch in the execution path depending on the state of a given CPU flag. Here, the instruction provides only a relative offset, which is added to the contents of the program counter (PC) as it points to the immediate next instruction. The relative offset is a signed single byte value in two's complement encoding (giving a range of -128...+127), which allows for branching up to half a page forwards and backwards.

On the one hand, this makes these instructions compact, fast and

Generally, an assembler will take care of this and we only have to provide the target address, not having to worry about relative addressing.

PC	Mnemonic	Instruction	Target
\$1000	BEQ \$1005	F0 03	
\$1002	PC pointing to next instruction		PC: \$1005

Diagram illustrating PC-relative addressing. The instruction at PC \$1000 is BEQ \$1005. The instruction format shows the opcode F0 and the offset 03. The PC is updated to \$1002 (PC pointing to next instruction). The target address is calculated as PC (\$1002) + Offset (03) = PC: \$1005.

```
BEQ $1005 ....branch.to location "$1005", if the zero flag is set.
               if the current address is $1000, this will give an offset of $03.
BCS $08C4 ....branch.to location "$08C4", if the carry flag is set.
               if the current address is $08D4, this will give an offset of $EE (-$12).
BCC $084A ....branch.to location "$084A", if the carry flag is clear.
```

MOS Technology, 1975



ADC Add Memory to Accumulator with Carry

addressing	assembler	opc	bytes	cycles
immediate	ADC #oper	69	2	2
zeropage	ADC oper	65	2	3
zeropage,X	ADC oper,X	75	2	4
absolute	ADC oper	6D	3	4
absolute,X	ADC oper,X	7D	3	4*
absolute,Y	ADC oper,Y	79	3	4*
(indirect,X)	ADC (oper,X)	61	2	6
(indirect),Y	ADC (oper),Y	71	2	5*

addressing	assembler	opc	bytes	cycles
immediate	AND #oper	29	2	2
zeropage	AND oper	25	2	3
zeropage,X	AND oper,X	35	2	4
absolute	AND oper	2D	3	4
absolute,X	AND oper,X	3D	3	4*
absolute,Y	AND oper,Y	39	3	4*
(indirect,X)	AND (oper,X)	21	2	6
(indirect),Y	AND (oper),Y	31	2	5*

ASL Shift Left One Bit (Memory or Accumulator)

```
C <- [76543210] <- 0          N Z C I D V
                                + + + - - -
```

addressing	assembler	opc	bytes	cycles
accumulator	ASL A	0A	1	2
zeropage	ASL oper	06	2	5
zeropage,X	ASL oper,X	16	2	6
absolute	ASL oper	0E	3	6
absolute,X	ASL oper,X	1E	3	7

BCC Branch on Carry Clear

```
branch on C = 0                N Z C I D V
                                - - - - -
```

addressing	assembler	opc	bytes	cycles
relative	BCC oper	90	2	2**

BCS Branch on Carry Set

```
branch on C = 1                N Z C I D V
                                - - - - -
```

addressing	assembler	opc	bytes	cycles
relative	BCS oper	B0	2	2**

BEQ Branch on Result Zero

```
branch on Z = 1                N Z C I D V
                                - - - - -
```

addressing	assembler	opc	bytes	cycles
relative	BEQ oper	F0	2	2**

BIT Test Bits in Memory with Accumulator

bits 7 and 6 of operand are transfered to bit 7 and 6 of SR (N,V);
the zero-flag is set according to the result of the operand AND
the accumulator (set, if the result is zero, unset otherwise).
This allows a quick check of a few bits at once without affecting
any of the registers, other than the status register (SR).

```
A AND M -> Z, M7 -> N, M6 -> V    N Z C I D V
                                M7 + - - - M6
```

addressing	assembler	opc	bytes	cycles
zeropage	BIT oper	24	2	3
absolute	BIT oper	2C	3	4

BMI Branch on Result Minus

```
branch on N = 1                N Z C I D V
                                - - - - -
```

addressing	assembler	opc	bytes	cycles
relative	BMI oper	30	2	2**

BNE Branch on Result not Zero

```
branch on Z = 0                N Z C I D V
                                - - - - -
```

addressing	assembler	opc	bytes	cycles
relative	BNE oper	D0	2	2**

BPL Branch on Result Plus

```
branch on N = 0                N Z C I D V
                                - - - - -
```

addressing	assembler	opc	bytes	cycles
relative	BPL oper	10	2	2**

BRK Force Break

BRK initiates a software interrupt similar to a hardware

interrupt,	N	Z	C	I	D	V
push PC+2, push SR	-	-	-	1	-	-

```

branch on V = 0
N Z C I D V
- - - - -

```

```

branch on V = 1
N Z C I D V
- - - - -

```

0	->	C		N	Z	C	I	D	V
				-	-	0	-	-	-

0	->	D	N	Z	C	I	D	V
			-	-	-	-	0	-

```
0 -> I      N Z C I D V
            - - - 0 - -
```

```
0 -> V      N Z C I D V
             - - - - - 0
```

A - M	N Z C I D V
	+ + + - - -

CPX Compare Memory and Index X

N	Z	C	I	D	V
+	+	+	-	-	-

addressing	assembler	opc	bytes	cycles
immediate	CPX #oper	E0	2	2
zeropage	CPX oper	E4	2	3
absolute	CPX oper	EC	3	4

CPY Compare Memory and Index Y

N	Z	C	I	D	V
+	+	+	-	-	-

addressing	assembler	opc	bytes	cycles
immediate	CPY #oper	C0	2	2
zeropage	CPY oper	C4	2	3
absolute	CPY oper	CC	3	4

DEC Decrement Memory by One

N	Z	C	I	D	V
+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
zeropage	DEC oper	C6	2	5
zeropage,X	DEC oper,X	D6	2	6
absolute	DEC oper	CE	3	6
absolute,X	DEC oper,X	DE	3	7

DEX Decrement Index X by One

N	Z	C	I	D	V
+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	DEX	CA	1	2

DEY Decrement Index Y by One

N	Z	C	I	D	V
+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	DEY	88	1	2

EOR Exclusive-OR Memory with Accumulator

N Z C I D V
+ + - - - -

addressing	assembler	opc	bytes	cycles
immediate	EOR #oper	49	2	2
zeropage	EOR oper	45	2	3
zeropage,X	EOR oper,X	55	2	4
absolute	EOR oper	4D	3	4
absolute,X	EOR oper,X	5D	3	4*
absolute,Y	EOR oper,Y	59	3	4*
(indirect,X)	EOR (oper,X)	41	2	6
(indirect),Y	EOR (oper),Y	51	2	5*

INC Increment Memory by One

N	Z	C	I	D	V
+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
zeropage	INC oper	E6	2	5
zeropage,X	INC oper,X	F6	2	6
absolute	INC oper	EE	3	6
absolute,X	INC oper,X	FE	3	7

INX Increment Index X by One

N Z C I D V
+ + - - - -

addressing	assembler	opc	bytes	cycles
implied	INX	E8	1	2

Y + 1 -> Y	N	Z	C	I	D	V
	+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	INYN	C8	1	2

JMP Jump to New Location

```
operand 1st byte -> PCL      N Z C I D V
operand 2nd byte -> PCH      - - - - - -
```

addressing	assembler	opc	bytes	cycles
absolute	JMP oper	4C	3	3
indirect	JMP (oper)	6C	3	5

JSR Jump to New Location Saving Return Address

```

push (PC+2),           N Z C I D V
operand 1st byte -> PCL  - - - - -
operand 2nd byte -> PCH

```

addressing	assembler	opc	bytes	cycles
absolute	JSR oper	20	3	6

LDA Load Accumulator with Memory

M	->	A		N	Z	C	I	D	V
				+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
immediate	LDA #oper	A9	2	2
zeropage	LDA oper	A5	2	3
zeropage,X	LDA oper,X	B5	2	4
absolute	LDA oper	AD	3	4
absolute,X	LDA oper,X	BD	3	4*
absolute,Y	LDA oper,Y	B9	3	4*
(indirect,X)	LDA (oper,X)	A1	2	6
(indirect),Y	LDA (oper),Y	B1	2	5*

LDX Load Index X with Memory

M	->	X	N	Z	C	I	D	V
			+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
immediate	LDX #oper	A2	2	2
zeropage	LDX oper	A6	2	3
zeropage,Y	LDX oper,Y	B6	2	4
absolute	LDX oper	AE	3	4
absolute,Y	LDX oper,Y	BE	3	4*

LDY Load Index Y with Memory

M	->	Y	N	Z	C	I	D	V
			+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
immediate	LDY #oper	A0	2	2
zeropage	LDY oper	A4	2	3
zeropage,X	LDY oper,X	B4	2	4
absolute	LDY oper	AC	3	4
absolute,X	LDY oper,X	BC	3	4*

LSR Shift One Bit Right (Memory or Accumulator)

0 -> [76543210] -> C

	N	Z	C	I	D	V
0	+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
accumulator	LSR A	4A	1	2
zeropage	LSR oper	46	2	5
zeropage,X	LSR oper,X	56	2	6
absolute	LSR oper	4E	3	6
absolute,X	LSR oper,X	5E	3	7

NOP No Operation

--- N Z C I D V - - - - -				
addressing	assembler	opc	bytes	cycles
implied	NOP	EA	1	2

ORA OR Memory with Accumulator

A OR M -> A N Z C I D V
 + + - - -

addressing	assembler	opc	bytes	cycles
immediate	ORA #oper	09	2	2
zeropage	ORA oper	05	2	3
zeropage,X	ORA oper,X	15	2	4
absolute	ORA oper	0D	3	4
absolute,X	ORA oper,X	1D	3	4*
absolute,Y	ORA oper,Y	19	3	4*
(indirect,X)	ORA (oper,X)	01	2	6
(indirect),Y	ORA (oper),Y	11	2	5*

PHA Push Accumulator on Stack

push A N Z C I D V
 - - - - -

addressing	assembler	opc	bytes	cycles
implied	PHA	48	1	3

PHP Push Processor Status on Stack

The status register will be pushed with the break flag and bit 5 set to 1.

push SR N Z C I D V
 - - - - -

addressing	assembler	opc	bytes	cycles
implied	PHP	08	1	3

PLA Pull Accumulator from Stack

pull A N Z C I D V
 + + - - -

addressing	assembler	opc	bytes	cycles
implied	PLA	68	1	4

PLP Pull Processor Status from Stack

The status register will be pulled with the break flag and bit 5 ignored.

pull SR N Z C I D V
 from stack

addressing	assembler	opc	bytes	cycles
implied	PLP	28	1	4

ROL Rotate One Bit Left (Memory or Accumulator)

C <- [76543210] <- C N Z C I D V
 + + + - - -

addressing	assembler	opc	bytes	cycles
accumulator	ROL A	2A	1	2
zeropage	ROL oper	26	2	5
zeropage,X	ROL oper,X	36	2	6
absolute	ROL oper	2E	3	6
absolute,X	ROL oper,X	3E	3	7

ROR Rotate One Bit Right (Memory or Accumulator)

C -> [76543210] -> C N Z C I D V
 + + + - - -

addressing	assembler	opc	bytes	cycles
accumulator	ROR A	6A	1	2
zeropage	ROR oper	66	2	5

zeropage,X	ROR oper,X	76	2	6
absolute	ROR oper	6E	3	6
absolute,X	ROR oper,X	7E	3	7

RTI Return from Interrupt

The status register is pulled with the break flag and bit 5 ignored. Then PC is pulled from the stack.

pull SR, pull PC N Z C I D V
 from stack

addressing	assembler	opc	bytes	cycles
implied	RTI	40	1	6

RTS Return from Subroutine

pull PC, PC+1 -> PC N Z C I D V
 - - - - -

addressing	assembler	opc	bytes	cycles
implied	RTS	60	1	6

SBC Subtract Memory from Accumulator with Borrow

A - M - C⁻ -> A N Z C I D V
 + + + - +

addressing	assembler	opc	bytes	cycles
immediate	SBC #oper	E9	2	2
zeropage	SBC oper	E5	2	3
zeropage,X	SBC oper,X	F5	2	4
absolute	SBC oper	ED	3	4
absolute,X	SBC oper,X	FD	3	4*
absolute,Y	SBC oper,Y	F9	3	4*
(indirect,X)	SBC (oper,X)	E1	2	6
(indirect),Y	SBC (oper),Y	F1	2	5*

SEC Set Carry Flag

1 -> C N Z C I D V
 - - 1 - - -

addressing	assembler	opc	bytes	cycles
implied	SEC	38	1	2

SED Set Decimal Flag

1 -> D N Z C I D V
 - - - 1 -

addressing	assembler	opc	bytes	cycles
implied	SED	F8	1	2

SEI Set Interrupt Disable Status

1 -> I N Z C I D V
 - - - 1 - -

addressing	assembler	opc	bytes	cycles
implied	SEI	78	1	2

STA Store Accumulator in Memory

A -> M N Z C I D V
 - - - - -

addressing	assembler	opc	bytes	cycles
zeropage	STA oper	85	2	3
zeropage,X	STA oper,X	95	2	4
absolute	STA oper	8D	3	4
absolute,X	STA oper,X	9D	3	5
absolute,Y	STA oper,Y	99	3	5
(indirect,X)	STA (oper,X)	81	2	6
(indirect),Y	STA (oper),Y	91	2	6

STX Store Index X in Memory

N Z C I D V
- - - - -

addressing	assembler	opc	bytes	cycles
zeropage	STX oper	86	2	3
zeropage,Y	STX oper,Y	96	2	4
absolute	STX oper	8E	3	4

STY Sore Index Y in Memory

Y -> M N Z C I D V
- - - - -

addressing	assembler	opc	bytes	cycles
zeropage	STY oper	84	2	3
zeropage,X	STY oper,X	94	2	4
absolute	STY oper	8C	3	4

TAX Transfer Accumulator to Index X

A	->	X	N	Z	C	I	D	V
			+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	TAX	AA	1	2

TAY Transfer Accumulator to Index Y

A -> Y	N Z C I D V
	+ + - - - -

addressing	assembler	opc	bytes	cycles
implied	TAY	A8	1	2

TSX Transfer Stack Pointer to Index X

SP -> X		N	Z	C	I	D	V
		+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	TSX	BA	1	2

TXA Transfer Index X to Accumulator

X	->	A	N	Z	C	I	D	V
			+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	TXA	8A	1	2

TXS Transfer Index X to Stack Register

X -> SP	N	Z	C	I	D	V
	-	-	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	TXS	9A	1	2

TYA Transfer Index Y to Accumulator

Y -> A		N	Z	C	I	D	V
		+	+	-	-	-	-

addressing	assembler	opc	bytes	cycles
implied	TYA	98	1	2

```
*   add 1 to cycles if page boundary is crossed
**  add 1 to cycles if branch occurs on same page
    add 2 to cycles if branch occurs to different page
```

```
Legend to Flags:  + .... modified
                  - .... not modified
                  1 .... set
                  0 .... cleared
                  M6 .... memory bit 6
                  M7 .... memory bit 7
```

Note on assembler syntax:
Some assemblers employ "OPC *oper" or a ".b" extension to the mnemonic for forced zeropage addressing.

Note on Read-Modify-Write instructions (*NMOS 6502 only*):
Some instructions like EOR, ASL, ROL, DEC, INC, etc., fetch a value from memory to modify it and to write the modified value back to the originating address. The original NMOS 6502 switches immediately into write mode after the read of the value, resulting in the unmodified value being written back to the address (while the value is modified in the next cycle), before the modified value is finally written to the destination. This may cause issues when writing to devices attached to the address bus that may trigger some action on this intermediate write operation.
This does not apply to the CMOS variants of the 6502.

Honorable Mention: Rev. A 6502 ROR, Pre-June 1976

Famously, the Rev. A 6502 as delivered from September 1975 to June 1976 had a "ROR bug". However, the "ROR" instruction isn't only missing from the original documentation, as it turns out, the chip is actually [missing crucial control lines](#), which would have been required to make this instruction work. The instruction is simply not implemented and it wasn't even part of the design. (This was actually added on popular demand in Rev. B, as rumor has it, demand by Steve Wozniak. Even, if not true, this makes for a good story. And how could there be a page on the 6502 without mentioning "Woz" once?) So, for all means, "ROR" is an undocumented or "illegal" instruction on the Rev. A 6502.

And this is how ROR behaves on these Rev. A chips, much like ASL: it shifts all bits to the left, shifting in a zero bit at the LSB side, but, unlike ASL, it does not shift the high-bit into the carry. (So there are no connections to the carry at all.)

```
ROR Rev. A (pre-June 1976)

As ASL, but does not update the carry.

N and Z flags are set correctly for the operation
performed.

[76543210] <- 0

                N Z C I D V
                + + - - - -

addressing  assembler  opc  bytes  cycles
-----
accumulator ROR A      6A   1      2
zeropage    ROR oper    66   2      5
zeropage,X  ROR oper,X  76   2      6
absolute    ROR oper    6E   3      6
absolute,X  ROR oper,X  7E   3      7
```

Compare Instructions

The 6502 MPU features three basic compare instructions in various address modes:

Instruction	Comparison
CMP	Accumulator and operand
CPX	X register and operand
CPY	Y register and operand

The various compare instructions subtract the operand from the respective register (as if the carry was set) without setting the result and adjust the N, Z, and C flags accordingly to this operation.
Flags will be set as follows:

Relation	Z	C	N
register < operand	0	0	<i>sign-bit of result</i>
register = operand	1	1	0
register > operand	0	1	<i>sign-bit of result</i>

And for the derivative relations "*less/greater than or equal*":

Relation	Z	C	N
register \leq operand	1	0	<i>sign-bit of result</i>
register \geq operand	1	1	<i>sign-bit of result</i>

Mind that the negative flag is not significant and all conditions may be evaluated by checking the carry and/or zero flag(s).

The BIT Instruction

The [BIT](#) instruction may be the most obscure instruction of the 6502: While other instructions serve a very clear purpose, like transferring values or performing basic arithmetic or logical operations, this one serves a rather specialized purpose, but it does so in a very general way. This purpose is bit testing.

Generally, testing of a particular bit is achieved by masking (isolating) this bit (or multiple bits) by an AND operation and then checking the zero flag (Z) by a BNE or BEQ instruction. This, however, destroys the contents of the accumulator. This is, where the BIT instruction comes in: much like the comparisons perform a subtraction without setting the result, the BIT instruction performs a logical AND without setting the result, but still reflects the result in the state of the zero flag (Z). Which allows for the same checks using the BNE or BEQ instructions, without affecting the contents of the accumulator.

Since the sign-bit is often used as a flag, testing this is also covered by the BIT instruction, which additionally to setting the zero flag also transfers bits 7 and 6 of the operand into the corresponding bits of the status register – which happen to be the negative (N) and overflow (V) flags. Therefore, bits 7 and 6 of the operand may be tested independently using the BMI/BPL and BVS/BVC instructions.

```

accumulator      operand
[76543210]  AND  [76543210]  == 0?
               ↓↓           ↓
               NV           Z

```

A Primer of 6502 Arithmetic Operations

The 6502 processor features two basic arithmetic instructions, ADC, *ADD with Carry*, and SBC, *SUBtract with Carry*. As the names suggest, these provide addition and subtraction for single byte operands and results. However, operations are not limited to a single byte range, which is where the carry flag comes in, providing the means for a single-bit carry (or borrow), to combine operations over several bytes.

In order to accomplish this, the carry is included in each of these operations: for additions, it is added (much like another operand); for subtractions, which are just an addition using the inverse of the operand (complement value of the operand), the role of the carry is inverted, as well. Therefore, it is crucial to set up the carry appropriately: for additions, the carry has to be initially cleared (using CLC), while for subtractions, it must be initially set (using SEC – more on SBC below).

```

                ;ADC: A = A + M + C
CLC             ;clear carry in preparation
LDA #2          ;load 2 into the accumulator
ADD #3          ;add 3 -> now 5 in accumulator

                ;SBC: A = A - M - C- ("C-": "not carry")
SEC            ;set carry in preparation
LDA #15         ;load 15 into the accumulator
SBC #8          ;subtract 8 -> now 7 in accumulator

```

Note: Here, we used immediate mode, indicated by the prefix "#" before the operand, to directly load a literal value. If there is no such "#" prefix, we generally mean to use the value stored at the address, which is given by the operand. As

we will see in the next example.)

To combine this for 16-bit values (2 bytes each), we simply chain the instructions for the next bytes to operate on, but this time without setting or clearing the carry.

Supposing the following locations for storing 16-bit values:

	low-byte	high-byte
first argument	\$1000	\$1001
second argument ...	\$1002	\$1003
result	\$1004	\$1005

we perform a 16-bit addition by:

```
CLC          ;prepare carry for addition
LDA $1000    ;load value at address $1000 into A (low byte of first argument)
ADC $1002    ;add low byte of second argument at $1002
STA $1004    ;store low byte of result at $1004
LDA $1001    ;load high byte of first argument
ADC $1003    ;add high byte of second argument
STA $1005    ;store high byte of result (result in $1004 and $1005)
```

and, conversely, for a 16-bit subtraction:

```
SEC          ;prepare carry for subtraction
LDA $1000    ;load value at address $1000 into A (low byte of first argument)
SBC $1002    ;subtract low byte of second argument at $1002
STA $1004    ;store low byte of result at $1004
LDA $1001    ;load high byte of first argument
SBC $1003    ;subtract high byte of second argument
STA $1005    ;store high byte of result (result in $1004 and $1005)
```

Note: Another, important preparatory step is to set the processor into binary mode by use of the CLD (*CLear Decimal flag*) instruction. (Compare the section on decimal mode below.) This has to be done only once.

Signed Values

Operations for unsigned and signed values are principally the same, the only difference being in how we interpret the values. Generally, the 6502 uses what is known as *two's complement* to represent negative values.

(In earlier computers, something known as ones' complement was used, where we simply flip all bits to their opposite state to represent a negative value. While simple, this came with a few drawbacks, like an additional value of negative zero, which are overcome by two's complement.)

In two's complement representation, we simply flip all the bits in a byte to their opposite (the same as an XOR by \$FF) and then add 1 to this.

E.g., to represent -4:

(We here use "\$" to indicate a hexadecimal number and "%" for binary notation. A dot is used to separate the high- and low-nibble, i.e. group of 4 bits.)

%0000.0100	4
XOR %1111.1111	255

%1111.1011	complement (all bits flipped)
+	1

%1111.1100	-4, two's complement

Thus, in a single byte, we may represent values in the range

from	-128	(%1000.0000 or \$80)
to	+127	(%0111.1111 or \$7F)

A notable feature is that the highest value bit (first bit from the left) will always be 1 for a negative value and always be 0 for a positive one, for which it is also known as the *sign bit*. Whenever we interpret a value as a signed number, a set sign bit indicates a negative value.

This works just the same for larger values, e.g., for a signed 16-bit value:

```
-512 = %1111.1110.0000.0000 = $FE $00
-516 = %1111.1101.1111.1100 = $FD $FC (mind how the +1 step carries over)
```

Notably, the binary operations are still the same as with unsigned values and provide the expected results:

<i>dec</i>	<i>binary</i>	<i>hex</i>	
100	%0110.0100	\$64	
+ -24	%1110.1000	\$E8	

76	%0100.1100	\$4C	(+ carry)

Note: We may now see how SBC actually works, by adding *ones' complement* of the operand to the accumulator. If we add 1 from the carry to the result, this effectively results in a subtraction in *two's complement* (the inverse of the operand + 1). If the carry happens to be zero, the result falls short by 1 in terms of two's complement, which is equivalent to adding 1 to the operand before the subtraction. Thus, the carry either provides the correction required for a valid two's complement representation or, if missing, results in a subtraction including a binary borrow.

Flags with ADC and SBC

Besides the carry flag (C), which allows us to chain multi-byte operations, the CPU sets the following flags on the result of an arithmetic operation:

```
zero flag (Z) ..... set if the result is zero, else unset
negative flag (N) ... the N flag always reflects the sign bit of the result
overflow flag (V) ... indicates overflow in signed operations
```

The latter may require explanation: how is signed overflow different from the carry flag? The overflow flag is about a certain ambiguity of the sign bit and the negative flag in signed context: if operands are of the same sign, the case may occur, where the sign bit flips (as indicated by a change of the negative flag), while the result is still of the same sign. This condition is indicated by the overflow flag. Notably, such an overflow can never occur, when the operands are of opposite signs.

E.g., adding positive \$40 to positive \$40:

	<i>acc.</i> <i>hex</i>	<i>acc.</i> <i>binary</i>	<i>flags</i> NVDIZC
LDA #\$40	\$40	%0100.0000	000000
ADC #\$40	\$80	%1000.0000	110000

Here, the change of the sign bit is unrelated to the actual value in the accumulator, it is merely a consequence of carry propagation from bit 6 to bit 7, the sign bit. Since both operands are positive, the result must be positive, as well.

The overflow flag (V) is of interest in signed context only and has no meaning in unsigned context.

Decimal Mode (BCD)

Besides binary arithmetic, the 6502 processor supports a second mode, binary coded decimal (BCD), where each byte, rather than representing a range of 0...255, represents two decimal digits packed into a single byte. For this, a byte is thought divided into two sections of 4 bits, the high- and the low-nibble. Only values from 0...9 are used for each nibble and a byte can represent a range of a 2-digit decimal value only, as in 0...99.

E.g.,

<i>dec</i>	<i>binary</i>	<i>hex</i>
14	%0001.0100	\$14
98	%1001.1000	\$98

Mind how this intuitively translates to hexadecimal notation, where figures A...F are never used.

Whether or not the processor is in decimal mode is determined by the decimal flag (D). If it is set (using SED) the processor will use BCD arithmetic. If it is cleared (using CLD), the processor is in binary mode. Decimal mode only affects instructions ADC and SBC (but not INC or DEC.)

Examples:

```
SED
CLC
LDA #$12
ADC #$44 ;accumulator now holds $56
```

```
SED
CLC
LDA #$28
ADC #$14 ;accumulator now holds $42
```

Mind that BCD mode is always unsigned:

```
acc. NVDIZC
SED
SEC
LDA #0    $00 001011
SBC #1    $99 101000
```

The carry flag and the zero flag work in decimal mode as expected. The negative flag is set similar to binary mode (and of questionable value.) The overflow flag has no meaning in decimal mode.

Multi-byte operations are just as in decimal mode: We first prepare the carry and then chain operations of the individual bytes in increasing value order, starting with the lowest value pair.

(It may be important to note that Western Design Center (WDC) version of the processor, the 65C02, always clears the decimal flag when it enters an interrupt, while the original NMOS version of the 6502 does not.)

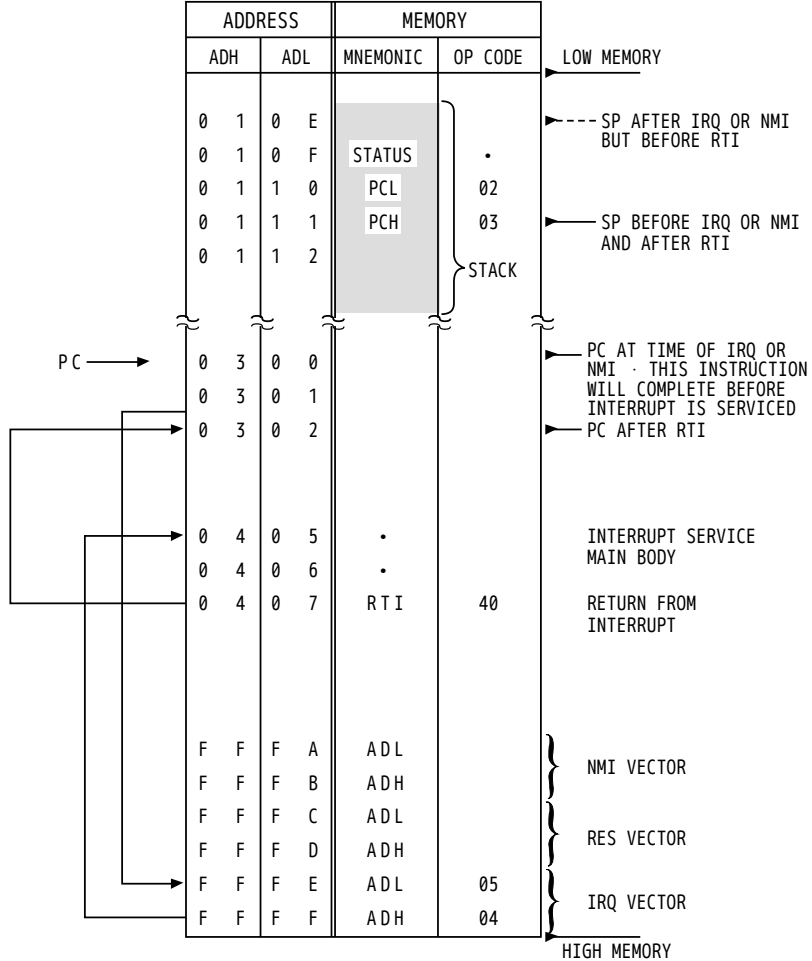
6502 Jump Vectors and Stack Operations

The 256 bytes processor stack of the 6502 is located at \$0100 ... \$01FF in memory, growing down from top to bottom.

There are three 2-byte address locations at the very top end of the 64K address space serving as jump vectors for reset/startup and interrupt operations:

```
$FFFA, $FFFB ... NMI (Non-Maskable Interrupt) vector
$FFFC, $FFFD ... RES (Reset) vector
$FFFE, $FFFF ... IRQ (Interrupt Request) vector
```


As an interrupt occurs, any instruction currently processed is completed first. Only then, the value of the program counter (PC) is put in high-low order onto the stack, followed by the value currently in the status register, and control will be transferred to the address location found in the respective interrupt vector. The registers stored on the stack are recovered at the end of an interrupt routine, as control is transferred back to the interrupted code by the RTI instruction.



IRQ, NMI, RTI, BRK OPERATION

(Reset after: MCS6502 Instruction Set Summary, MOS Technology, Inc.)

Similarly, as a JSR instruction is encountered, PC is dumped onto the stack and recovered by the JSR instruction. (Here, the value stored is actually the address *before* the location, the program will eventually return to. Thus, the effective return address is PC+1.)



JSR, RTS OPERATION

(Reset after: MCS6502 Instruction Set Summary, MOS Technology, Inc.)

Curious Interrupt Behavior

- If the instruction was a taken branch instruction with 3 cycles execution time (without crossing page boundaries), the interrupt will trigger only after an extra CPU cycle.
- On the NMOS6502, an NMI hardware interrupt occurring at the start of a BRK instruction will hijack the BRK instruction, meaning, the BRK instruction will be executed as normal, but the NMI vector will be used instead of the IRQ vector.
- The 65C02 will clear the decimal flag on any interrupts (and BRK).

The Break Flag and the Stack

Interrupts and stack operations involving the status register (or P register) are the only instances, the break flag appears (namely on the stack). It has no representation in the CPU and can't be accessed by any instruction.

- The break flag will be set to on (1), whenever the transfer was caused by software (BRK or PHP).
- The break flag will be set to zero (0), whenever the transfer was caused by a hardware interrupt.
- The break flag will be masked and cleared (0), whenever transferred from the stack to the status register, either by PLP or during a return from interrupt (RTI).

Therefore, it's somewhat difficult to inspect the break flag in order to discern a software interrupt (BRK) from a hardware interrupt (NMI or IRQ) and the mechanism is seldom used. Accessing a break mark put in the extra byte following a BRK instruction is even more cumbersome and probably involves indexed zeropage operations.

Bit 5 (unused) of the status register will be set to 1, whenever the register is pushed to the stack. Bits 5 and 4 will always be ignored, when transferred to the status register.

E.g.,

1)

```
SR: N V - B D I Z C
    0 0 - - 0 0 1 1
```

```
PHP -> 0 0 1 1 0 0 1 1 = $33
```

```
PLP <- 0 0 - - 0 0 1 1 = $03
```

but:

```
PLA <- 0 0 1 1 0 0 1 1 = $33
```

2)

```
LDA #$32 ;00110010
```

```
PHA -> 0 0 1 1 0 0 1 0 = $32
```

```
PLP <- 0 0 - - 0 0 1 0 = $02
```

3)

```
LDA #$C0
```

```
PHA -> 1 1 0 0 0 0 0 0 = $C0
```

```
LDA #$08
```

```
PHA -> 0 0 0 0 1 0 0 0 = $08
```

```
LDA #$12
```

```
PHA -> 0 0 0 1 0 0 1 0 = $12
```

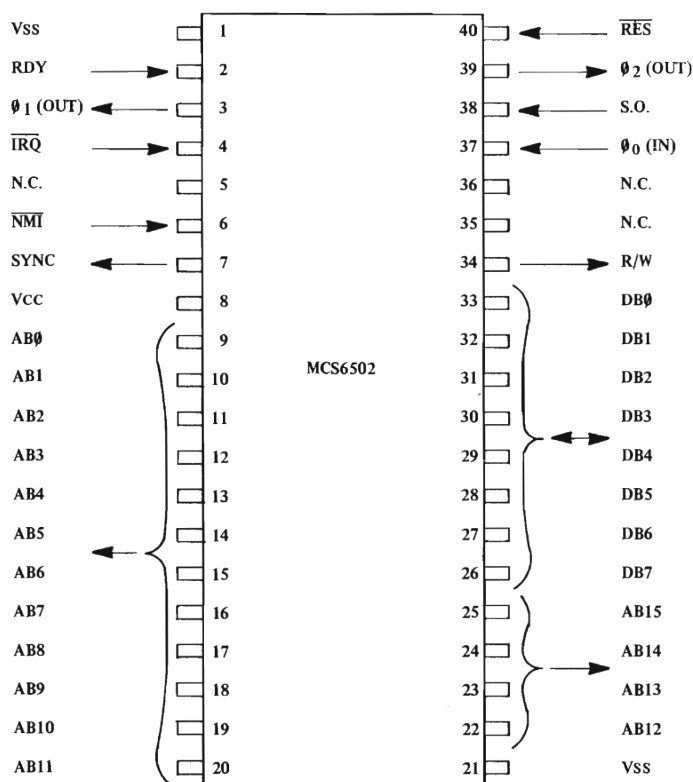
```
RTI
```

```
SR: 0 0 - - 0 0 1 0 = $02
```

```
PC: $C008
```

Mind that most emulators are displaying the status register (SR or P) in the state as it would be currently pushed to the stack, with bits 4 and 5 on, adding a bias of \$30 to the register value. Here, we chose to rather omit this virtual presence of these bits, since there isn't really a slot for them in the hardware.

Pinout (NMOS 6502)



Vcc, Vss	supply voltage (Vcc: +5 V DC ± 5%, Vss: max. +7 V DC)
Φ0...2	clock
AB0-AB15	address bus
DB0-DB7	data bus
R/W	read/write
RDY	ready
S.O.	set overflow (future I/O interface)
SYNC	sync (goes high on opcode fetch phase)
I ⁻ R ⁻ Q ⁻	interrupt request (active low)
N ⁻ M ⁻ I ⁻	non maskable interrupt (active low)
R ⁻ E ⁻ S ⁻	reset (active low)
N.C.	no connection

Source: MCS6500 Microcomputer Family Hardware Manual. MOS Technology, Inc., 1976.

The 65xx-Family:

Type	Features, Comments
6502	NMOS, 16 bit address bus, 8 bit data bus
6502A	accelerated version of 6502
6502C	accelerated version of 6502, additional halt pin, CMOS
65C02	WDC version, additional instructions and address modes, up to 14MHz
6503, 6505, 6506	12 bit address bus [4 KiB]
6504	13 bit address bus [8 KiB], no NMI
6507	13 bit address bus [8 KiB], no interrupt lines
6509	20 bit address bus [1 MiB] by bankswitching
6510	as 6502 with additional 6 bit I/O-port
6511	integrated micro controler with I/O-port, serial interface, and RAM (Rockwell)
65F11	as 6511, integrated FORTH interpreter
7501	as 6502, HMOS
8500	as 6510, CMOS
8502	as 6510 with switchable 2 MHz option, 7 bit I/O-port
65816 (65C816)	16 bit registers and ALU, 24 bit address bus [16 MiB], up to 24 MHz (Western Design Center)
65802 (65C802)	as 65816, pin compatible to 6502, 64 KiB address bus, up to 16 MHz